



An overview of SQL

Dec 2017



© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License
See <http://creativecommons.org/licenses/by-nc-sa/4.0/>



3 Adams E-mail: SCOTTADAMS@AOL.COM


1997 © 1996 United Feature Syndicate, Inc.(NYC)

Structured Query Language

- Pronounced “Sequel”
 - Originally called Sequel but changed for trademark reasons
- Dates to 1974
 - Written by IBM (Chamberlin and Boyce)
 - Based on “A Relational Model of Data for Large Shared Data Banks” by Edward Codd
 - First commercialised by Oracle
 - Standardised in 1986



Relational terminology (from Codd)



Relational Database

- Every row in a table has the same attributes (columns)
 - Relations are either tables or views on those tables
- A primary key for each row uniquely identifies it
- A foreign key points to another table's primary key



Relational database

Id	Firstname	Lastname	birthdate
2587	John	Hopkins	5/12/1973
7789	Henry	Gleeson	1/5/1985
22398	Eleanor	Richardson	10/6/1996

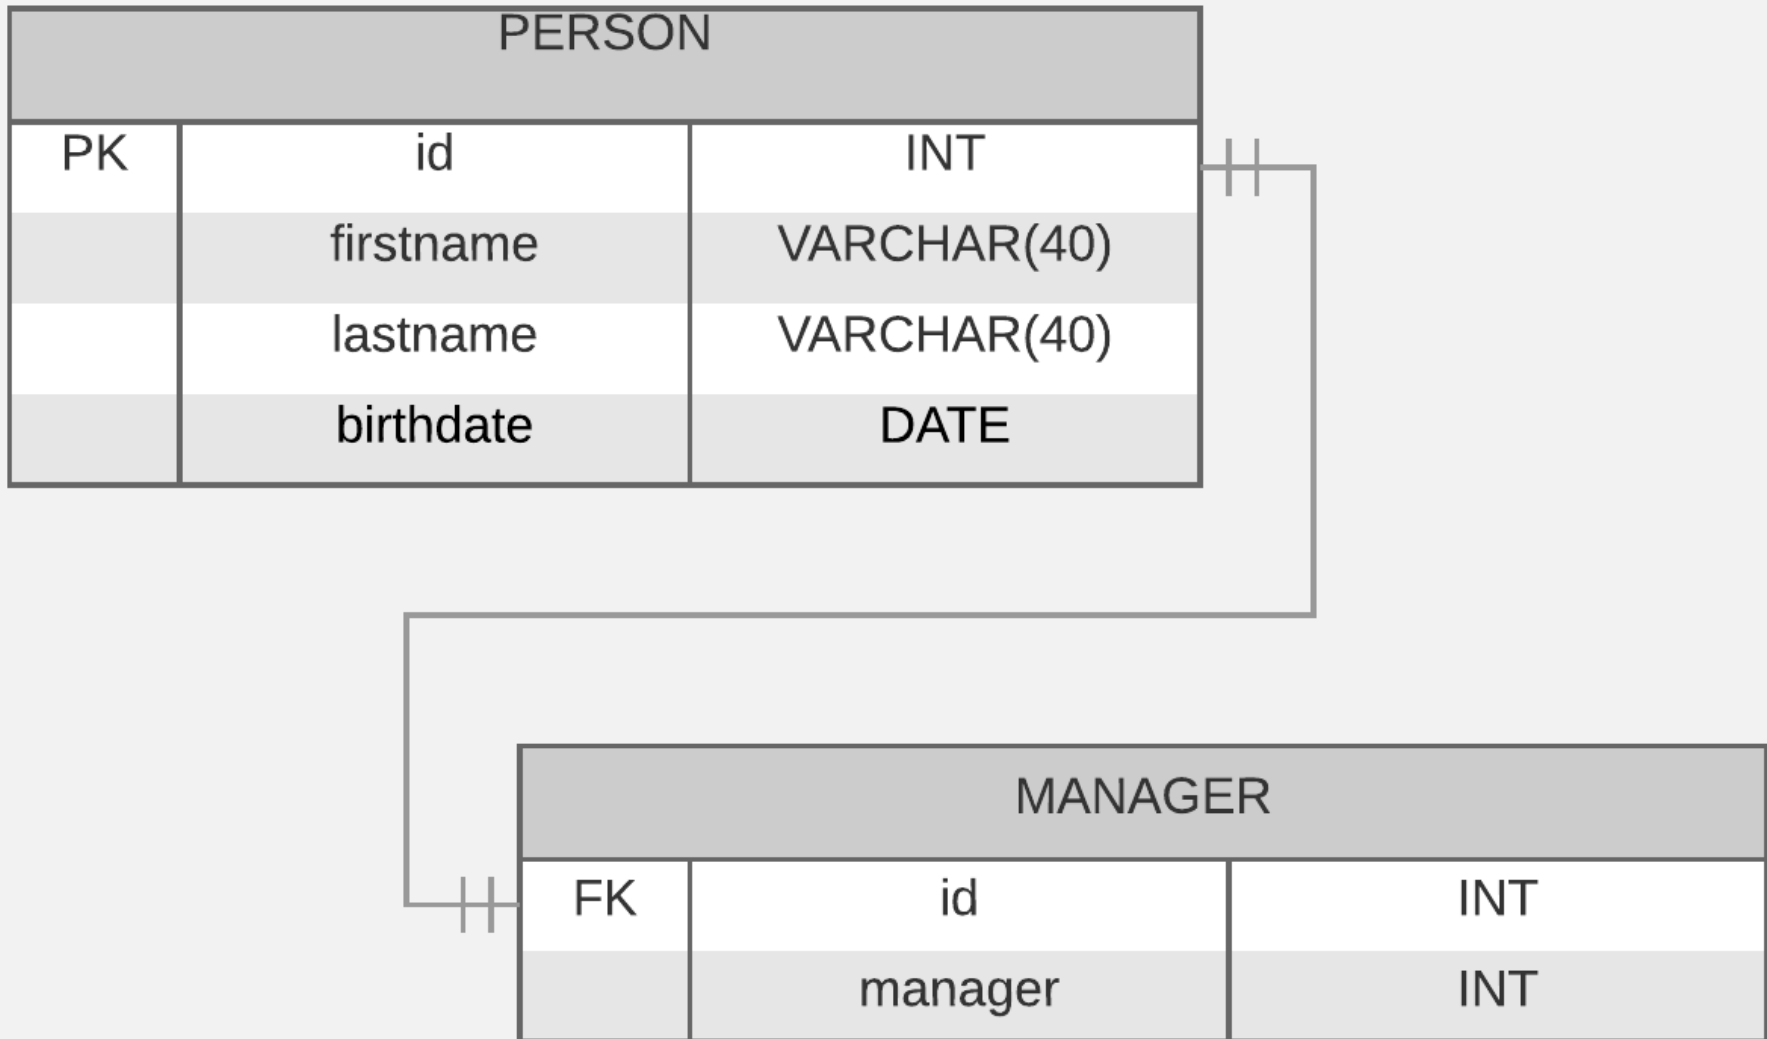


Why are we looking at SQL today?

- SQL and variations are widely used
 - Not just for relational databases
- Hive / SparkSQL
 - SQL over big data using map-reduce techniques
- Siddhi / KSQL / StreamingSQL
 - SQL queries over real-time streaming data
- Other SQL interfaces
 - e.g. SQL into Sloan Digital Sky Survey



Entity Relationship Diagram



SQL STATEMENTS

corresponding to the previous diagram

```
CREATE TABLE `PERSON` (  
  `id` INT,  
  `firstname` VARCHAR(40),  
  `lastname` VARCHAR(40),  
  `birthdate` DATE,  
  PRIMARY KEY (`id`)  
);
```

```
CREATE TABLE `MANAGER` (  
  `id` INT,  
  `manager` INT,  
  KEY `FK` (`id`)  
);
```



INSERT

```
INSERT INTO person  
  (id, firstname, lastname, birthdate)  
values  
  (564, "Henry", "Gleeson", "1968-12-5");
```

```
INSERT INTO person  
  (id, firstname, lastname, birthdate)  
values  
  (2343, "Eleanor", "Smith", "1995-1-9");
```



SELECT

SELECT * FROM person;

id	firstname	lastname	birthdate
-----	-----	-----	-----
564	Henry	Gleeson	1968-12-5
2343	Eleanor	Smith	1995-1-9



SELECT

SELECT * FROM person WHERE id = 564;

id	firstname	lastname	birthdate
-----	-----	-----	-----
564	Henry	Gleeson	1968-12-5



SELECT

**SELECT * FROM person WHERE
firstname = "Eleanor";**

id	firstname	lastname	birthdate
-----	-----	-----	-----
2343	Eleanor	Smith	1995-1-9



SELECT

SELECT firstname, lastname **FROM** person
ORDER BY firstname;

firstname	lastname
-----	-----
Eleanor	Smith
Henry	Gleeson



SELECT

```
SELECT firstname, lastname FROM person  
ORDER BY lastname LIMIT 1;
```

firstname	lastname
-----	-----
Henry	Gleeson



SELECT

```
SELECT AVG(birthdate) FROM person;
```

```
AVG(birthdate)
```

```
-----
```

```
1981.5
```



Functions

- MIN
- MAX
- AVG
- COUNT
- SUM



OTHER COMMANDS

- DELETE

DELETE FROM person WHERE ID=564;

- UPDATE


UPDATE PERSON

SET firstname = Henrietta WHERE ID=564;



This is a very brief introduction!

- We will learn more from the exercises
- There are lots of resources on the Web



Apache Hive

<http://hive.apache.org>



- Just like SQL except it generates Map Reduce jobs
- Works on Hadoop and Spark
 - Embedded into Spark as SparkSQL
- Includes DDL (Data Definition Language) as well as SQL
- Makes many processing tasks very simple



Hive example

```
CREATE TABLE page_view(viewTime INT, userid BIGINT,  
                        page_url STRING, referrer_url STRING,  
                        ip STRING COMMENT 'IP Address of the User')  
COMMENT 'This is the page view table'  
PARTITIONED BY(dt STRING, country STRING)  
STORED AS SEQUENCEFILE;
```

```
LOAD DATA LOCAL INPATH /tmp/pv_2008-06-08_us.txt INTO TABLE page_view  
PARTITION(date='2008-06-08', country='US')
```

```
INSERT OVERWRITE TABLE xyz_com_page_views  
SELECT page_views.*  
FROM page_views  
WHERE page_views.date >= '2008-03-01' AND page_views.date <=  
'2008-03-31' AND  
      page_views.referrer_url like '%xyz.com';
```



SparkSQL

- Integrates into existing Spark programs
 - Mixes SQL with Python, Scala or Java
- Integrates data from CSV, Avro, Parquet, JDBC, ODBC, JSON, etc
 - Including joins across them
- Fully supports Apache Hive
 - *If you build it with Hive support*
- Fits into the resilient scalable model of Spark



Spark SQL example

```
from pyspark.sql import SQLContext, Row
sqlContext = SQLContext(sc)
```

```
lines = sc.textFile("examples/src/main/resources/people.txt")
parts = lines.map(lambda l: l.split(","))
people = parts.map(lambda p: Row(name=p[0], age=int(p[1])))
```

```
schemaPeople = sqlContext.createDataFrame(people)
schemaPeople.registerTempTable("people")
```

```
teenagers = sqlContext.sql("SELECT name FROM people WHERE age >= 13
AND age <= 19")
```

```
teenNames = teenagers.map(lambda p: "Name: " + p.name)
for teenName in teenNames.collect():
    print(teenName)
```



DataFrame

Based on Python and R dataframes

- Column based object used by SQL
- Offers SQL like programming
- Supports algebraic optimisation and code gen
- E.g. in Scala:

```
means = users.where(users["age"] > 20)
               .groupBy("city")
               .avg("income")
```

And they run up to 2-5x faster than equivalent computations expressed via the functional API.

More SQL

df.

```
select('postcode','id').  
withColumn('first_pc',  
  split(df.postcode, '\s'[0]).  
  where((col("first_pc") == 'SW11') or  
        (col("first_pc") == 'OX1'))).  
groupBy('first_pc').  
agg({"id": "count"}).show()
```



User Defined Functions

- In SQL a User Defined Function is an extension that helps



Questions?



© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License
See <http://creativecommons.org/licenses/by-nc-sa/4.0/>