

Map-Reduce and Scaling Big Data Processing

 © Paul Fremont 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <http://creativecommons.org/licenses/by-nc-sa/4.0/>

Contents

- Understanding Map Reduce
 - Functional programming patterns applied for scalability
- Hadoop
 - Map-reduce in Hadoop
 - Python
 - Java
 - HDFS
 - Yarn
 - Pig and Hive
- Further reading

 © Paul Fremont 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <http://creativecommons.org/licenses/by-nc-sa/4.0/>

Original 2008 Google Paper

MapReduce: Simplified Data Processing on Large Clusters

Jeffrey Dean and Sanjay Ghemawat


jeff@google.com, sanjay@google.com

Google, Inc.

Abstract

MapReduce is a programming model and an associated implementation for processing and generating large data sets. Users specify a *map* function that processes a key/value pair to generate a set of intermediate key/value pairs, and a *reduce* function that merges all intermediate values associated with the same intermediate key. Many real world tasks are expressible in this model, as shown

given day, etc. Most such computations are conceptually straightforward. However, the input data is usually large and the computations have to be distributed across hundreds or thousands of machines in order to finish in a reasonable amount of time. The issues of how to parallelize the computation, distribute the data, and handle failures conspire to obscure the original simple computation with large amounts of complex code to deal with these issues.


 © Paul Fremont 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <http://creativecommons.org/licenses/by-nc-sa/4.0/>

Yahoo 2007

Map-Reduce-Merge: Simplified Relational Data Processing on Large Clusters

Hung-chih Yang, Ali Dasdan
Yahoo!
Sunnyvale, CA, USA
(hcyang,dasdan)@yahoo-inc.com

Ruey-Lung Hsiao, D. Stott Parker
Computer Science Department, UCLA
Los Angeles, CA, USA
(rlhsiao,stott)@cs.ucla.edu

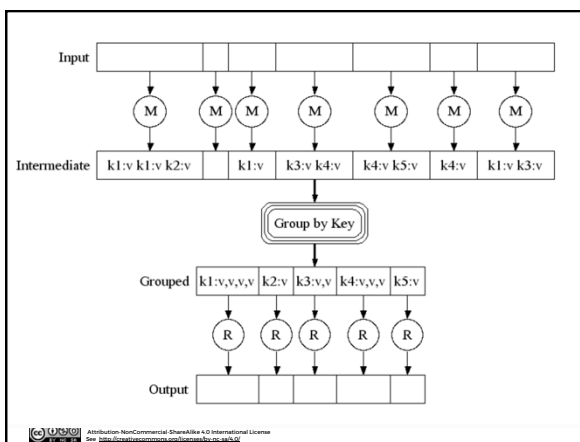
 © Paul Fremont 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <http://creativecommons.org/licenses/by-nc-sa/4.0/>

Class Exercise

- Find a small piece of paper and write your university and day/month of birth on **one**.
- You don't need the year

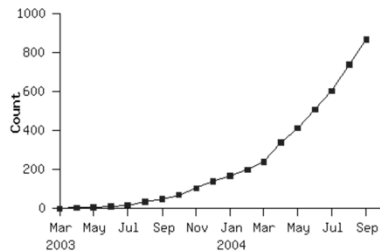
Portsmouth
5 October

 © Paul Fremont 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <http://creativecommons.org/licenses/by-nc-sa/4.0/>



Google's early use of MR

Map Reduce programs in their code repository



© Paul Fremont 2005. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <http://creativecommons.org/licenses/by-nc-sa/4.0/>

Map Reduce example in words

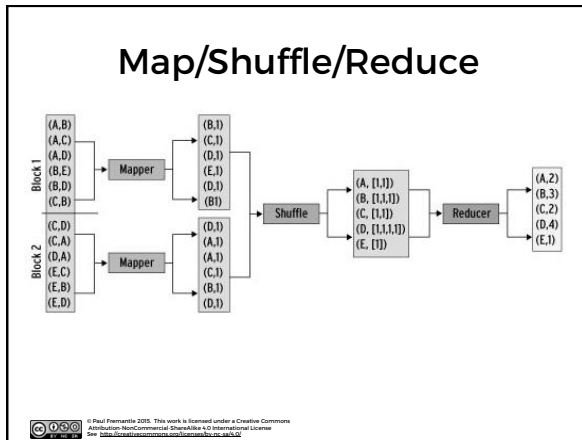
- Do a word count on 1000 books:
 - First count each book (**map** wc onto book)
 - Then **reduce** the outputs to a global wordcount across all books

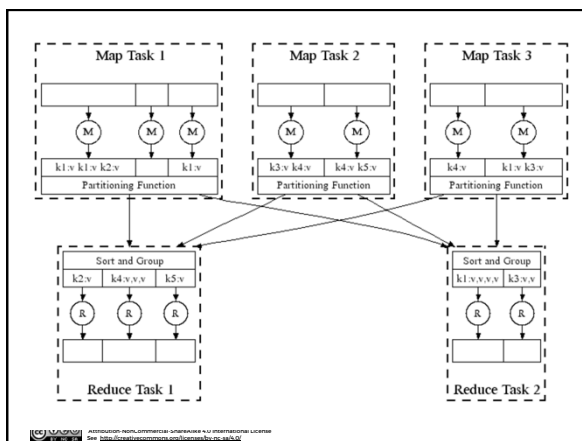
© Paul Fremont 2005. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <http://creativecommons.org/licenses/by-nc-sa/4.0/>

Efficiency

- Reduce phase:
 - We can theoretically process each word in parallel
- How?
 - Shuffle / Sort the results from the map phase by key (word)
 - Partition by keys
 - Parallelize the reduce phase

© Paul Fremont 2005. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <http://creativecommons.org/licenses/by-nc-sa/4.0/>






Map Reduce in Real Life

- Analysing web logs
 - Summarise by user / cookie
 - Then aggregate to identify who did what
- Analysing twitter data
 - Who retweeted
 - Who was retweeted the most
- Almost all big data problems can be re-factored into Map Reduce
 - Some more efficiently than others

© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

Tuning

- **Fault tolerance**
 - Simply re-execute work that fails
- **Performance:**
 - Partitioning the data
 - Moving the work to near the data

 © Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <http://creativecommons.org/licenses/by-nc-sa/4.0/>

Apache Hadoop

- **The most famous and popular Map Reduce framework**
 - **Open Source**
 - Written in Java, but supports other languages
 - Runs Map Reduce workloads across a cloud or cluster of machines
 - Supports a distributed filesystem to store data for these jobs
 - Provides reliability when servers in the cluster fail

 © Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <http://creativecommons.org/licenses/by-nc-sa/4.0/>

Components of Hadoop



 © Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <http://creativecommons.org/licenses/by-nc-sa/4.0/>

Summary

- Understanding the Map Reduce Model
- How is it implemented in Hadoop
- HDFS
- Yarn

 © Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

Questions?

 © Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <https://creativecommons.org/licenses/by-nc-sa/4.0/>
