

An overview of SQL

Dec 2017


 © Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <https://creativecommons.org/licenses/by-nc-sa/4.0/>




 © Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

Structured Query Language

- Pronounced “Sequel”
 - Originally called Sequel but changed for trademark reasons
- Dates to 1974
 - Written by IBM (Chamberlin and Boyce)
 - Based on “A Relational Model of Data for Large Shared Data Banks” by Edward Codd
 - First commercialised by Oracle
 - Standardised in 1986

 © Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

Relational terminology (from Codd)



© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

Relational Database

- Every row in a table has the same attributes (columns)
 - Relations are either tables or views on those tables
- A primary key for each row uniquely identifies it
- A foreign key points to another table's primary key

© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

Relational database

Id	Firstname	Lastname	birthdate
2587	John	Hopkins	5/12/1973
7789	Henry	Gleeson	1/5/1985
22398	Eleanor	Richardson	10/6/1996

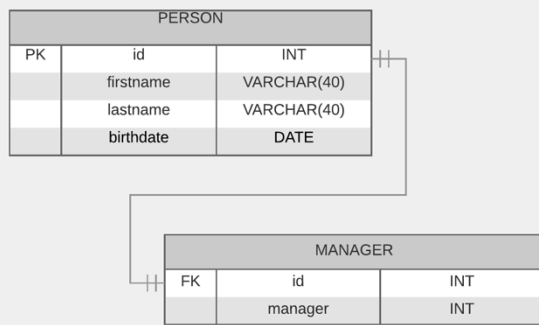
© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

Why are we looking at SQL today?

- SQL and variations are widely used
 - Not just for relational databases
- Hive / SparkSQL
 - SQL over big data using map-reduce techniques
- Siddhi / KSQL / StreamingSQL
 - SQL queries over real-time streaming data
- Other SQL interfaces
 - e.g. SQL into Sloan Digital Sky Survey

 © Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

Entity Relationship Diagram



SQL STATEMENTS corresponding to the previous diagram

```

CREATE TABLE 'PERSON' (
  'id' INT,
  'firstname' VARCHAR(40),
  'lastname' VARCHAR(40),
  'birthdate' DATE,
  PRIMARY KEY ('id')
);

CREATE TABLE 'MANAGER' (
  'id' INT,
  'manager' INT,
  KEY 'FK' ('id')
);
  
```

 © Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

INSERT

```
INSERT INTO person
(id, firstname, lastname, birthdate)
values
(564, "Henry", "Gleeson", "1968-12-5");
INSERT INTO person
(id, firstname, lastname, birthdate)
values
(2343, "Eleanor", "Smith", "1995-1-9");
```

 © Paul Fremantle 2015. This work is licensed under a Creative Commons
Attribution-NonCommercial-ShareAlike 4.0 International License
See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

SELECT

```
SELECT * FROM person;
```

id	firstname	lastname	birthdate
564	Henry	Gleeson	1968-12-5
2343	Eleanor	Smith	1995-1-9

 © Paul Fremantle 2015. This work is licensed under a Creative Commons
Attribution-NonCommercial-ShareAlike 4.0 International License
See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

SELECT

```
SELECT * FROM person WHERE id = 564;
```

id	firstname	lastname	birthdate
564	Henry	Gleeson	1968-12-5

 © Paul Fremantle 2015. This work is licensed under a Creative Commons
Attribution-NonCommercial-ShareAlike 4.0 International License
See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

SELECT

**SELECT * FROM person WHERE
firstname = "Eleanor";**

id	firstname	lastname	birthdate
2343	Eleanor	Smith	1995-1-9

 © Paul Fremantle 2015. This work is licensed under a Creative Commons
Attribution-NonCommercial-ShareAlike 4.0 International License
See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

SELECT

**SELECT firstname, lastname FROM person
ORDER BY firstname;**

firstname	lastname
Eleanor	Smith
Henry	Gleeson

 © Paul Fremantle 2015. This work is licensed under a Creative Commons
Attribution-NonCommercial-ShareAlike 4.0 International License
See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

SELECT

**SELECT firstname, lastname FROM person
ORDER BY lastname LIMIT 1;**

firstname	lastname
Henry	Gleeson

 © Paul Fremantle 2015. This work is licensed under a Creative Commons
Attribution-NonCommercial-ShareAlike 4.0 International License
See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

SELECT

```
SELECT AVG(birthdate) FROM person;
```

```
AVG(birthdate)
```

```
-----  
1981.5
```

 © Paul Fremont 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

Functions

- MIN
- MAX
- AVG
- COUNT
- SUM

 © Paul Fremont 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

OTHER COMMANDS


- DELETE

```
DELETE FROM person WHERE ID=564;
```

- UPDATE


```
UPDATE PERSON
```

```
SET firstname = Henrietta WHERE ID=564;
```

 © Paul Fremont 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

This is a very brief introduction!

- We will learn more from the exercises
- There are lots of resources on the Web



© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <http://creativecommons.org/licenses/by-nc-sa/4.0/>

Apache Hive

<http://hive.apache.org>



- Just like SQL except it generates Map Reduce jobs
- Works on Hadoop and Spark
 - Embedded into Spark as SparkSQL
- Includes DDL (Data Definition Language) as well as SQL
- Makes many processing tasks very simple

© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <http://creativecommons.org/licenses/by-nc-sa/4.0/>

Hive example

```
CREATE TABLE page_view(viewTime INT, userid BIGINT,
  page_url STRING, referrer_url STRING,
  ip STRING COMMENT 'IP Address of the User')
COMMENT 'This is the page view table'
PARTITIONED BY(dt STRING, country STRING)
STORED AS SEQUENCEFILE;

LOAD DATA LOCAL INPATH '/tmp/pv_2008-06-08_us.txt' INTO TABLE page_view
PARTITION(date='2008-06-08', country='US');

INSERT OVERWRITE TABLE xyz_com_page_views
SELECT page_views.*
FROM page_views
WHERE page_views.date >= '2008-03-01' AND page_views.date <=
'2008-03-31' AND
  page_views.referrer_url like '%xyz.com';
```

© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <http://creativecommons.org/licenses/by-nc-sa/4.0/>

SparkSQL

- Integrates into existing Spark programs
 - Mixes SQL with Python, Scala or Java
- Integrates data from CSV, Avro, Parquet, JDBC, ODBC, JSON, etc
 - Including joins across them
- Fully supports Apache Hive
 - *If you build it with Hive support*
- Fits into the resilient scalable model of Spark

 © Paul Fremont 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

Spark SQL example

```
from pyspark.sql import SQLContext, Row
sqlContext = SQLContext(sc)

lines = sc.textFile("examples/src/main/resources/people.txt")
parts = lines.map(lambda l: l.split(", "))
people = parts.map(lambda p: Row(name=p[0], age=int(p[1])))

schemaPeople = sqlContext.createDataFrame(people)
schemaPeople.registerTempTable("people")

teenagers = sqlContext.sql("SELECT name FROM people WHERE age >= 13
AND age <= 19")

teenNames = teenagers.map(lambda p: "Name: " + p.name)
for teenName in teenNames.collect():
    print(teenName)
```


 © Paul Fremont 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

DataFrame

Based on Python and R dataframes

- Column based object used by SQL
- Offers SQL like programming
- Supports algebraic optimisation and code gen
- E.g. in Scala:

```
means = users.where(users["age"] > 20)
               .groupBy("city")
               .avg("income")
```

 And they run up to 2-5x faster than equivalent computations expressed via the functional API.

More SQL

```
df.  
select('postcode','id').  
withColumn('first_pc',  
split(df.postcode, '\s'[0]).  
where((col("first_pc") == 'SW11') or  
      (col("first_pc") == 'OX1'))).  
groupBy('first_pc').  
agg({"id": "count"}).show()
```

 © Paul Fremantle 2015. This work is licensed under a Creative Commons
Attribution-NonCommercial-ShareAlike 4.0 International License
See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

User Defined Functions

- In SQL a User Defined Function is an extension that helps

 © Paul Fremantle 2015. This work is licensed under a Creative Commons
Attribution-NonCommercial-ShareAlike 4.0 International License
See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

Questions?

 © Paul Fremantle 2015. This work is licensed under a Creative Commons
Attribution-NonCommercial-ShareAlike 4.0 International License
See <https://creativecommons.org/licenses/by-nc-sa/4.0/>
