


## Big Data Engineering

### NoSQL databases

© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. For more information see <http://creativecommons.org/licenses/by-nc-sa/4.0/>

---

---

---

---


---

---

---

## Contents

- Why NoSQL?
- ReCAP
- BigTable and Dynamo
- A summary of a few NoSQL databases
  - MongoDB, Cassandra, Couchbase,

© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. For more information see <http://creativecommons.org/licenses/by-nc-sa/4.0/>

---

---

---

---


---

---

---

## Why NoSQL?

- Availability
  - Need better scaling capabilities
  - Elasticity
- Different schema approaches
  - Graphs, Key Values, Document, Sparse Columns, etc
- More appropriate balance in read/write performance
- Better integration with REST/SOA/Cloud

© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. For more information see <http://creativecommons.org/licenses/by-nc-sa/4.0/>

---

---

---

---

---

---

---

## NoSQL history

- Not just a recent thing ☺
- IBM IMS (Information Management System)
  - Launched in 1968
  - Used to store the bill of materials for the Saturn V rocket
  - Hierarchical model
- Still in widespread use today



© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. For more information, see <http://creativecommons.org/licenses/by-nc-sa/4.0/>

---

---

---

---

---

---


---

---

## ReCAP

- You can have 2 out of three:

- Consistent
  - ACID
- Available
  - HA / Accessible 24x7
- Partitioned
  - Able to split into different datacentres
  - Survive network down



© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. For more information, see <http://creativecommons.org/licenses/by-nc-sa/4.0/>

---

---

---

---

---

---

---

---

## NoSQL parents

- Amazon Dynamo
  - Eventually consistent
- Google BigTable
  - Supporting very large rows
- LDM
  - Graph database



© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. For more information, see <http://creativecommons.org/licenses/by-nc-sa/4.0/>

---

---

---

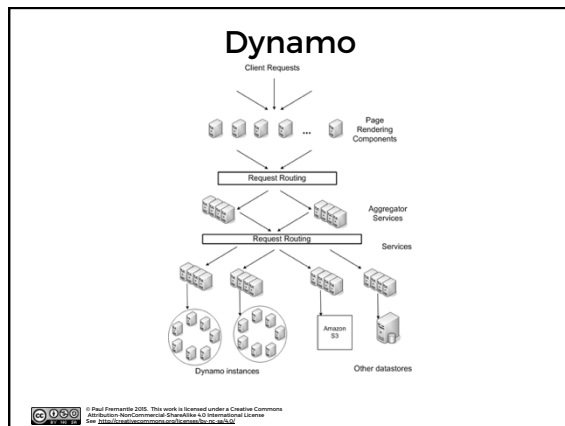
---

---

---

---

---




---

---

---

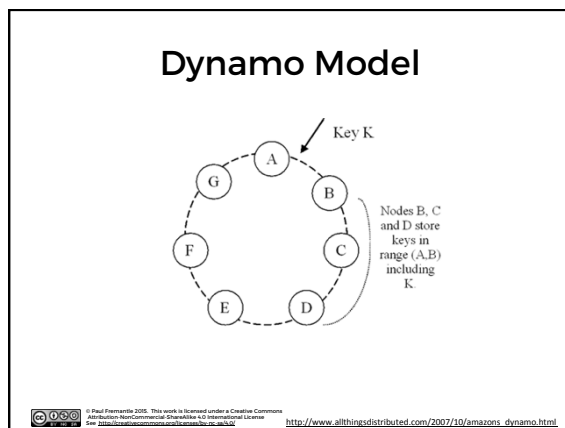
---

---

---

---

---




---

---

---

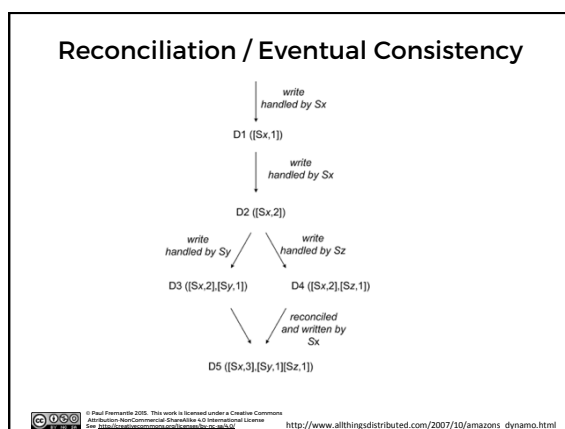
---

---

---

---

---




---

---

---

---

---


---

---

---

## Dynamo Techniques


Problem	Technique	Advantage
Partitioning	Consistent Hashing	Incremental Scalability
High Availability for writes	Vector clocks with reconciliation during reads	Version size is decoupled from update rates.
Handling temporary failures	Sloppy Quorum and hinted handoff	Provides high availability and durability guarantee when some of the replicas are not available.
Recovering from permanent failures	Anti-entropy using Merkle trees	Synchronizes divergent replicas in the background.
Membership and failure detection	Gossip-based membership protocol and failure detection.	Preserves symmetry and avoids having a centralized registry for storing membership and node liveness information.

 © Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. <https://creativecommons.org/licenses/by-nc-sa/4.0/>

## Google BigTable

- Optimized to support very large data
  - Not just many rows, but rows that cannot fit into the memory of a single server
  - Column Families allow each row to live across servers
- This table dates back to 2005

Project name	Table size (TB)	Compression ratio	# Cells (billions)	# Column Families	# Locality Groups	% in memory	Latency-sensitive?
<i>Crawl</i>	800	11%	1000	16	8	0%	No
<i>Crawl</i>	50	33%	200	2	2	0%	No
<i>Google Analytics</i>	20	29%	10	1	1	0%	Yes
<i>Google Analytics</i>	200	14%	80	1	1	0%	Yes
<i>Google Base</i>	2	31%	10	29	3	15%	Yes
<i>Google Earth</i>	0.5	64%	8	7	2	33%	Yes
<i>Google Earth</i>	70	—	9	8	3	0%	No
<i>Orkut</i>	9	—	0.9	8	5	1%	Yes
<i>Personalized Search</i>	4	47%	6	93	11	5%	Yes

 © Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. <https://creativecommons.org/licenses/by-nc-sa/4.0/>

## Current NoSQL Databases

- Too many to list!
- Popular databases include:
  - MongoDB
  - Couchbase
  - Apache Cassandra
  - Apache HBase
  - Voldemort
  - Redis
  - Riak
  - Etc, etc

 © Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. <https://creativecommons.org/licenses/by-nc-sa/4.0/>

## “NewSQL”

- ACID databases that aim to provide HA and Partition safety
  - VoltDB
  - NuoDB
  - Google Spanner
  - MemSQL
  - SAP HANA
- Also there are some backend engines for MySQL that aim to provide this:
  - MySQL Cluster
  - TokuDB

© Paul Fremantle 2016. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

---

---

---

---

---

---


---

---


## Vitess

<https://github.com/youtube/vitess>

Components



<https://freo.me/vitess-pres>




---

---

---

---

---

---

---

---

## In Memory Databases

- Memory is relatively much cheaper than it used to be
- Uses snapshots or transaction logs to ensure durability
- *Some NoSQL, some NewSQL*
  - SAP Hana
  - Redis
  - VoltDB
  - MemSQL
  - Apache Geode

© Paul Fremantle 2016. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See <https://creativecommons.org/licenses/by-nc-sa/4.0/>

---

---

---

---

---

---

---

---

## Top ten databases

283 systems in ranking, November 2015

Rank	Rank			DBMS	Database Model	Score		
	Nov 2015	Oct 2015	Nov 2014			Nov 2015	Oct 2015	Nov 2014
1.	1.	1.	1.	Oracle	Relational DBMS	1480.95	+13.99	+28.82
2.	2.	2.	2.	MySQL	Relational DBMS	1286.84	+7.88	+7.77
3.	3.	3.	3.	Microsoft SQL Server	Relational DBMS	1122.33	-0.90	-97.87
4.	4.	5.	4.	MongoDB	Document store	304.61	+11.34	+59.87
5.	5.	4.	5.	PostgreSQL	Relational DBMS	285.69	+3.56	+28.33
6.	6.	6.	6.	DB2	Relational DBMS	202.52	-4.28	-3.71
7.	7.	7.	7.	Microsoft Access	Relational DBMS	140.96	-0.87	+2.12
8.	8.	9.	8.	Cassandra	Wide column store	132.92	+3.91	+40.93
9.	9.	8.	9.	SQLite	Relational DBMS	103.45	+0.78	+8.17
10.	10.	11.	10.	Redis	Key-value store	102.41	+3.61	+20.06

<http://db-engines.com/en/ranking>


© 2015 DB-Engines. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. For more information, see <http://creativecommons.org/licenses/by-nc-sa/4.0/>

## Next 20

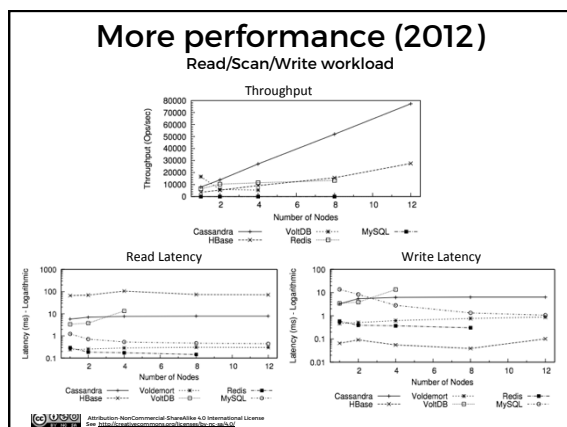
11.	11.	10.	11.	SAP Adaptive Server	Relational DBMS	83.71	-1.93	-0.91
12.	12.	12.	12.	Solr	Search engine	79.77	+0.71	+2.96
13.	13.	13.	13.	Teradata	Relational DBMS	77.08	+3.64	+9.85
14.	14.	16.	14.	Elasticsearch	Search engine	74.77	+4.55	+31.71
15.	15.	15.	15.	HBase	Wide column store	56.46	-0.78	+9.49
16.	16.	17.	16.	Hive	Relational DBMS	54.91	+1.35	+18.30
17.	17.	14.	17.	FileMaker	Relational DBMS	51.73	+1.95	+0.39
18.	18.	20.	18.	Splunk	Search engine	44.61	+1.11	+12.44
19.	19.	21.	19.	SAP HANA	Relational DBMS	39.62	+0.52	+11.26
20.	20.	18.	20.	Informix	Relational DBMS	38.46	+0.12	+2.88
21.	21.	23.	21.	Neo4j	Graph DBMS	34.04	+0.63	+9.39
22.	22.	19.	22.	Memcached	Key-value store	32.39	+0.82	-0.20
23.	23.	27.	23.	MariaDB	Relational DBMS	26.64	+2.01	+9.35
24.	24.	22.	24.	CouchDB	Document store	26.37	-0.49	+0.57
25.	25.	24.	25.	Couchbase	Document store	25.82	-0.37	+4.33
26.	26.	26.	26.	Firebird	Relational DBMS	22.56	-0.33	+4.95
27.	27.	25.	27.	Netezza	Relational DBMS	21.84	+0.60	+3.76
28.	28.	31.	28.	Amazon DynamoDB	Multi-model	21.75	+0.42	+9.43
29.	29.	28.	29.	Microsoft Azure SQL Database	Relational DBMS	19.46	+0.62	+4.85
30.	29.	29.	30.	Vertica	Relational DBMS	19.41	+0.45	+5.46

## Performance (2012)

50%/50% reads/writes



Rabl, Tilmann, et al. "Solving big data challenges for enterprise application performance management." Proceedings of the VLDB Endowment 5.12 (2012): 1724-1735.




---

---

---

---

---

---

---

---

### Summary of Performance benchmark (2012)

- **Cassandra** had the best throughput but high latency
- **Voldemort** had the best and most stable latency but lower throughput
- **HBase** had low performance per node but scaled well
  - Low write latency
- **Redis, MySQL and VoltDB** did not scale as well in multi-node setups

© 2012 Paul Fremantle. All rights reserved. Attribution-NonCommercial-ShareAlike 4.0 International License. See <http://creativecommons.org/licenses/by-nc-sa/4.0/>

---

---

---

---

---

---

---

---

### Key Value databases

- A persistent associative array or dictionary
- Simple access and fits well with programming models (especially MR)
- Indexing on other data is not often possible and can be slow

© 2012 Paul Fremantle. All rights reserved. Attribution-NonCommercial-ShareAlike 4.0 International License. See <http://creativecommons.org/licenses/by-nc-sa/4.0/>

---

---

---

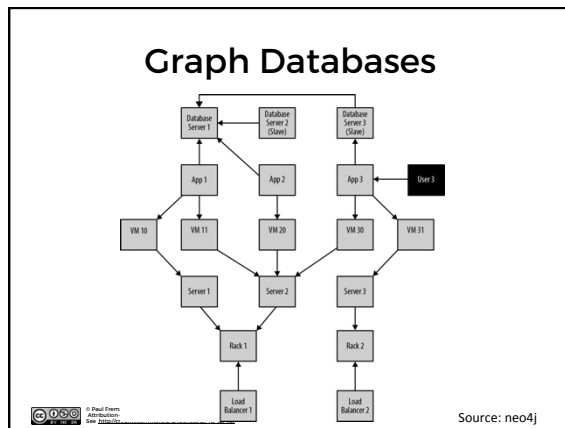
---

---

---

---

---




---

---

---

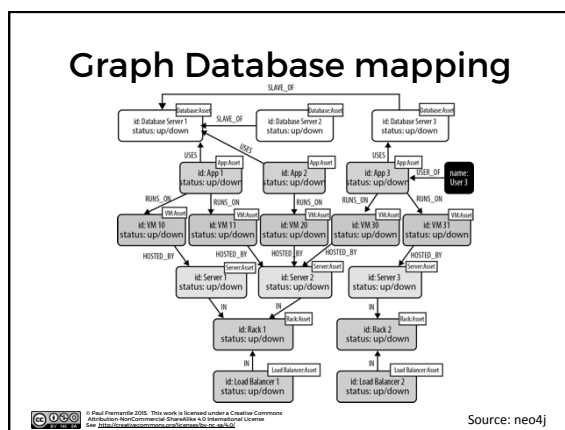
---

---

---

---

---




---

---

---

---

---

---

---

---

## Apache Cassandra

- **Masterless / Symmetric**
  - Every node is equal and you can write to any node as well as read
- **Shared Nothing architecture**
  - Each server has its own disk
- **Based on Dynamo**
  - for automatic sharding and eventual consistency
- **And BigTable**
  - For "Column Families"
- **Donated to Apache by Facebook**
  - Now mostly developed by DataStax

Source: DataStax

---

---

---

---

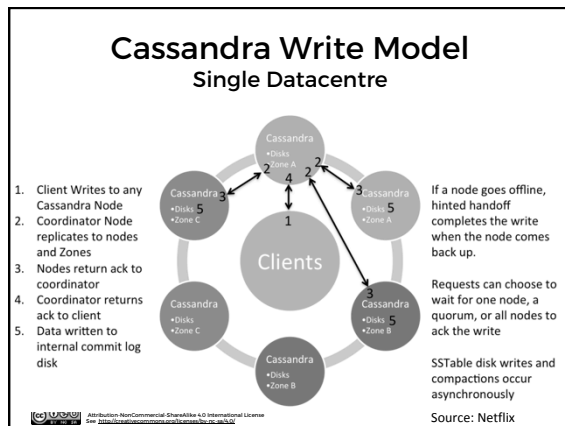
---

---

---

---






---

---

---

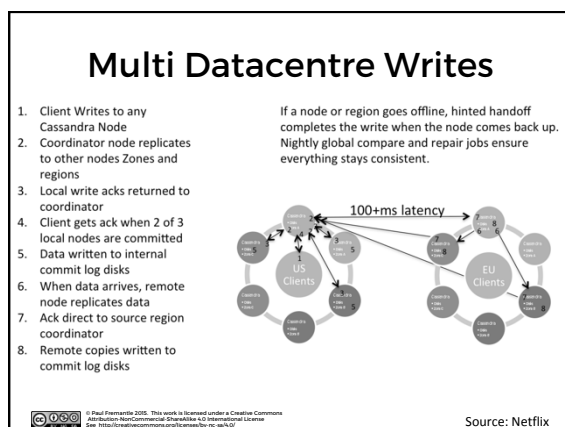
---

---

---

---

---




---

---

---


---

---

---

---

---




---

---

---

---

---


---

---

---

## The numbers

Per Node	48 Nodes	96 Nodes	144 Nodes	288 Nodes
Per Server Writes/s	10,900 w/s	11,460 w/s	11,900 w/s	11,456 w/s
Mean Server Latency	0.0117 ms	0.0134 ms	0.0148 ms	0.0139 ms
Mean CPU %Busy	74.4 %	75.4 %	72.5 %	81.5 %
Disk Read	5,600 KB/s	4,590 KB/s	4,060 KB/s	4,280 KB/s
Disk Write	12,800 KB/s	11,590 KB/s	10,380 KB/s	10,080 KB/s
Network Read	22,460 KB/s	23,610 KB/s	21,390 KB/s	23,640 KB/s
Network Write	18,600 KB/s	19,600 KB/s	17,810 KB/s	19,770 KB/s

 © Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. <https://creativecommons.org/licenses/by-nc-sa/4.0/>

Source: Netflix

## Cassandra Model

- **Keyspaces** are roughly equivalent to SQL Databases
  - Encapsulate replication strategies
- **Column Families** roughly equivalent to SQL tables
- Generally a different approach vs SQL
  - Writes are cheap
  - Indexes are expensive
  - Normalization is not the goal

 © Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. <https://creativecommons.org/licenses/by-nc-sa/4.0/>


## Cassandra Model cont.

- Inserts are the same as updates
  - No read first
- Data can be marked with a Time to Live (TTL)
  - Automatically deleted
- Deletes are not instant
  - Deleted rows are marked with a tombstone
  - Eventually cleaned up
  - Can re-appear if you do not run node repair after a node failure

 © Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. <https://creativecommons.org/licenses/by-nc-sa/4.0/>

## CQL

- A variant of SQL written specifically for Cassandra
  - The preferred model of access
  - Replaces the old “Thrift” API
- Attempts to have some compatibility with normal SQL
  - e.g. you can use either **KEYSPACE** or **TABLE** interchangeably

 If Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. For more information see <https://creativecommons.org/licenses/by-nc-sa/4.0/>

## CQL examples

**SELECT** name, occupation FROM users  
WHERE userid IN (199, 200, 207);

However, some queries are not permitted:

**SELECT** firstname, lastname FROM users WHERE  
birth\_year = 1981 AND country = 'FR';

Requires a large scan of the database and  
cannot give a predictable time response:

**ALLOW FILTERING** will make this run  
anyway

 If Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. For more information see <https://creativecommons.org/licenses/by-nc-sa/4.0/>

## INSERT / UPDATE

**INSERT INTO** NerdMovies (movie, director, main\_actor, year)  
VALUES ('Serenity', 'Joss Whedon', 'Nathan Fillion', 2005)  
*USING TTL 86400;*

- Every row can have a specified expiry time
- Inserts work even if the data is already there, unless you specify:


**INSERT INTO** NerdMovies (movie, director, main\_actor, year)  
VALUES ('Serenity', 'Joss Whedon', 'Nathan Fillion', 2005)  
*IF NOT EXISTS*  
*USING TTL 86400;*

This can have unpredictable timing because it requires read-before-write

 If Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. For more information see <https://creativecommons.org/licenses/by-nc-sa/4.0/>

## Non-SQL data types

- Sets
  - CREATE TABLE cycling.cyclist\_career\_teams ( id UUID PRIMARY KEY, lastname text, teams **set**<text> );
- Lists
  - CREATE TABLE cycling.upcoming\_calendar ( year int, month int, events **list**<text>, PRIMARY KEY ( year, month) );
- Maps
  - CREATE TABLE cycling.cyclist\_teams ( id UUID PRIMARY KEY, lastname text, firstname text, teams **map**<int,text> );
- Tuples
  - CREATE TABLE cycling.popular (rank int PRIMARY KEY, cinfo **tuple**<text,text,int> );

 © Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. For more information, see <http://creativecommons.org/licenses/by-nc-sa/4.0/>

---

---

---

---

---

---

---

---

## Direct support for JSON

```
INSERT INTO cycling.cyclist_category
JSON '{
  "category" : "GC",
  "points" : 780,
  "id" : "829aa84a-4bba-411f-
a4fb-38167a987cda",
  "lastname" : "SUTHERLAND" }';
```

 © Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. For more information, see <http://creativecommons.org/licenses/by-nc-sa/4.0/>

---

---

---

---

---

---

---

---

## cassandra.yaml

- Configuration of the major parts of the system
  - Datacentres, Racks, Cluster name
  - Authentication and Authorization
  - Partitioner
  - Data Storage location
  - Cacheing
  - Network topology and ports
  - Etc, etc

 © Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. For more information, see <http://creativecommons.org/licenses/by-nc-sa/4.0/>

---

---

---

---

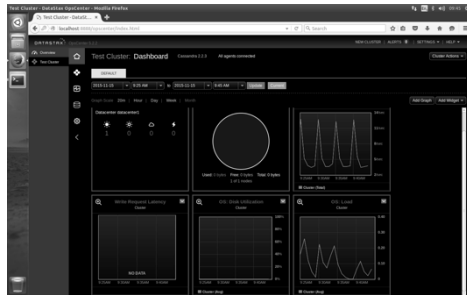
---

---

---

---

## DataStax OpsCenter



© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. <https://creativecommons.org/licenses/by-nc-sa/4.0/>

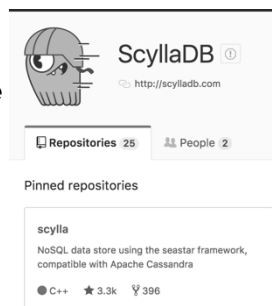
## OpsCenter

- Part of DataStax Cassandra distribution
  - Community edition has limited features
  - Enterprise edition expands these
- Not open source, but free to use in the community edition
  - Requires an agent on each Cassandra node
  - It will install this via SSH if possible

© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. <https://creativecommons.org/licenses/by-nc-sa/4.0/>


## ScyllaDB

- A C++ “clone” of Cassandra
- Also Open Source
- Claims to be significantly faster



© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. <https://creativecommons.org/licenses/by-nc-sa/4.0/>

## Scylla vs Cassandra



© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. For more information see <https://creativecommons.org/licenses/by-nc-sa/4.0/>

---

---

---

---

---

---

---

---

## Questions?

© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. For more information see <https://creativecommons.org/licenses/by-nc-sa/4.0/>

---

---

---

---

---

---

---

---