

## Big Data Engineering

### Realtime Big Data Processing

 © Paul Fremantle 2015. This work is licensed under a Creative Commons  
Attribution NonCommercial ShareAlike 4.0 International License  
For more information see <https://creativecommons.org/licenses/by-nc-sa/4.0/>

---

---

---

---

---

---

---

## Streaming

- Continuous data flow
  - “Unbounded streams of data”
- Usually uses a message distribution system
  - JMS, Apache Kafka, ZeroMQ, MQTT
- An unbounded set of events with time
  - $\langle t_1, E_1 \rangle, \langle t_2, E_2 \rangle, \dots, \langle t_n, E_n \rangle, \dots$

 © Paul Fremantle 2015. This work is licensed under a Creative Commons  
Attribution NonCommercial ShareAlike 4.0 International License  
For more information see <https://creativecommons.org/licenses/by-nc-sa/4.0/>

---

---

---

---

---

---

---

## Stream processing categorization

- Simple event processing
  - Working on an event at a time
    - e.g. filter out all events where the wind speed > 50 mph
- Event stream processing
  - Time-based processing of a single stream of events
    - Average wind speed over the last hour compared to the average over the last day
- Complex Event Processing
  - Correlation of events across different streams
    - Emergency calls correlated with wind speed in real time

 © Paul Fremantle 2015. This work is licensed under a Creative Commons  
Attribution NonCommercial ShareAlike 4.0 International License  
For more information see <https://creativecommons.org/licenses/by-nc-sa/4.0/>

---

---

---

---

---

---

---


## Comparing Databases with Real-Time systems

	Database Applications	Event-driven Applications
Query Paradigm	Ad-hoc queries or requests	Continuous standing queries
Latency	Seconds, hours, days	Milliseconds or less
Data Rate	Hundreds of events/sec	Tens of thousands of events/sec or more



© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. <http://creativecommons.org/licenses/by-nc-sa/4.0/>

## Lambda Architecture



© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. <http://creativecommons.org/licenses/by-nc-sa/4.0/>


## Approaches to Streaming

- **Pure streaming**
  - Each event is processed as it comes in
- **Micro-batch**
  - Small batches of events are processed
  - Typically trades flexibility for performance
- **Shared nothing**
  - You can process events on any system in the cluster
- **Stateful / Partitioned**
  - The event must be processed on a system that has the correct state in memory

© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. <http://creativecommons.org/licenses/by-nc-sa/4.0/>

## Data distribution

- You need to get the events to the processing systems

 © Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. <http://creativecommons.org/licenses/by-nc-sa/4.0/>

---

---

---

---


---

---

---

## MQTT

- Very simple, lightweight, fast
- No built in support for clustering / big-data
  - But can make up for it by being very fast
- Used a lot in IoT

 © Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. <http://creativecommons.org/licenses/by-nc-sa/4.0/>

---

---

---


---

---

---

---

## Apache Kafka



 © Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. <http://creativecommons.org/licenses/by-nc-sa/4.0/>

Source: <http://www.slideshare.net/charmalloc/>

---

---

---

---


---

---

---

## Kafka

- Many of the approaches we've seen:
  - Partitioning
  - Multiple brokers
  - Elastically scalable
  - Supports clusters of co-ordinated consumers
  - Automatic re-election of leaders

 © Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. For more information, see <https://creativecommons.org/licenses/by-nc-sa/4.0/>

---

---

---

---

---

---

---

## Kafka exactly-once semantics



**Mathias Verraes**  
@mathiasverraes

 Follow

There are only two hard problems in distributed systems: 2. Exactly-once delivery 1. Guaranteed order of messages 2. Exactly-once delivery

RETWEETS 6,775 LIKES 4,727 

10:40 AM - 14 Aug 2015

 69  6.8K  4.7K 

 © Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. For more information, see <https://creativecommons.org/licenses/by-nc-sa/4.0/>

---

---

---

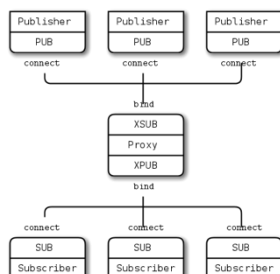
---


---

---

---

## ZeroMQ



 © Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. For more information, see <https://creativecommons.org/licenses/by-nc-sa/4.0/>

---

---

---

---

---

---

---

## Processing the data

 © Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. <http://creativecommons.org/licenses/by-nc-sa/4.0/>

---

---

---

---


---

---


---

---

## Apache Storm



Note: another DAG

 © Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. <http://creativecommons.org/licenses/by-nc-sa/4.0/>

---

---

---

---

---

---

---

---

## Apache Storm

- Originally developed by BackType
  - Nathan Marz
- Acquired by Twitter
- Open Sourced and then donated to Apache
- Became a top level project in 2014
  - <http://storm.apache.org>

 © Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. <http://creativecommons.org/licenses/by-nc-sa/4.0/>

---

---

---

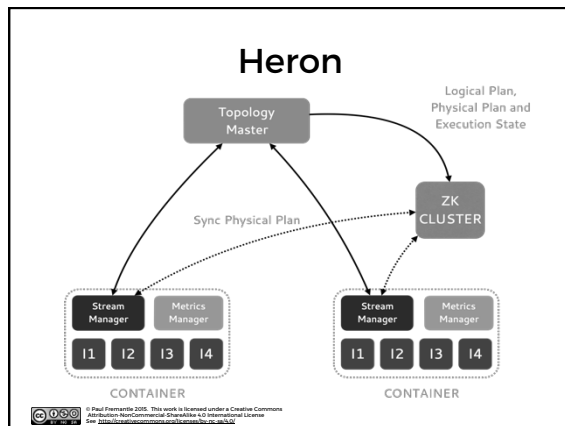
---

---

---

---

---




---

---

---

---

---

---

---

---

## Heron: Key Features

- Fully API compatible with Apache Storm
- Task isolation
- Developer productivity
- Ease of manageability
- Use of mainstream languages C++/Java/Python

© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. <https://creativecommons.org/licenses/by-nc-sa/4.0/>

17

---

---

---

---

---


---

---

---

## Heron

- In production at Twitter for >2 years
- Going into production at Microsoft, WeChat
- Donation to CNCF




---

---

---

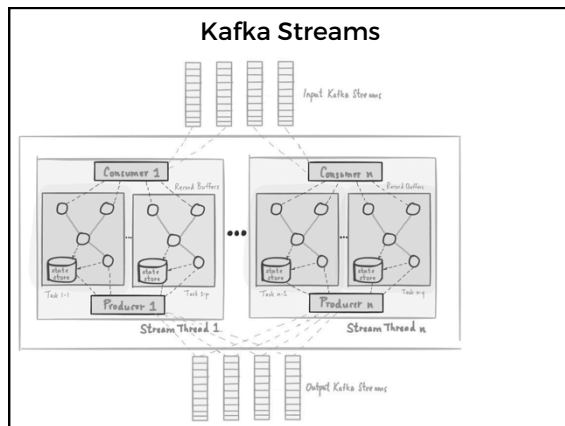
---

---

---

---

---




---

---

---

---

---

---

---

---

### Kafka Streams

- Event-at-a-time processing (not microbatch) with millisecond latency
- Stateful processing including distributed joins and aggregations
- A convenient DSL
- Windowing with out-of-order data using a DataFlow-like model
- Distributed processing and fault-tolerance with fast failover
- Reprocessing capabilities so you can recalculate output when your code changes
- No-downtime rolling deployments

© 2015 Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. <https://creativecommons.org/licenses/by-nc-sa/4.0/>

---

---

---

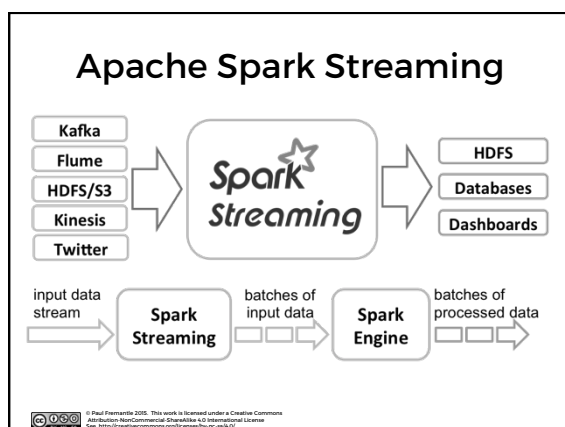
---

---

---

---

---




---

---

---

---

---


---

---

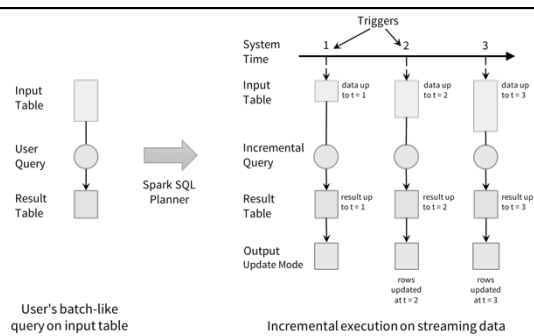
---

## Structured Streams in Spark

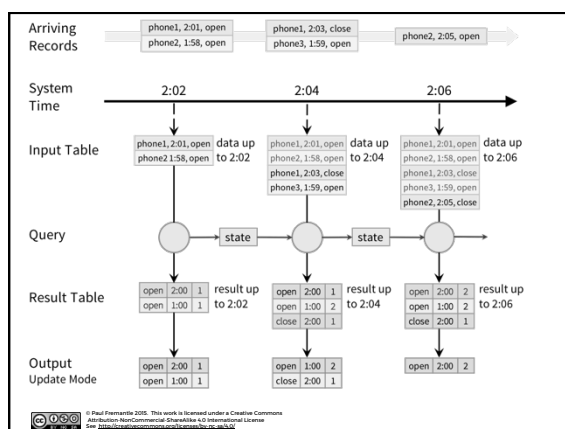
- Since Spark 2.0, there is a much better approach



© 2016 Paul Fremantle 2016. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. <https://creativecommons.org/licenses/by-nc-sa/4.0/>



© 2016 Paul Fremantle 2016. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. <https://creativecommons.org/licenses/by-nc-sa/4.0/>



© 2016 Paul Fremantle 2016. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. <https://creativecommons.org/licenses/by-nc-sa/4.0/>



## Siddhi

- A stateful query model
- SQL-like language for querying streams of data
  - Extended with **windows**
    - Time, Event count, batches
  - Partitioned
    - Based on data in the events
  - Pattern matching
    - A then B then C within window



© Paul Fremantle 2016. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. For more information, see <http://creativecommons.org/licenses/by-nc-sa/4.0/>

---

---

---

---

---


---

---

---

## Siddhi

- Apache Licensed Open Source on Github
  - <https://github.com/wso2/siddhi/>
- Pluggable into Storm, Spark and Kafka Streams
- Supports millions of events/sec
- [http://freo.me/DEBS\\_Siddhi](http://freo.me/DEBS_Siddhi)




---

---

---

---

---

---

---

---

## SiddhiQL

```
FROM login_stream#window.time(10 min)
SELECT ip,
       count(ip) as loginCount,
       cityId
GROUP BY ip
HAVING loginCount > 10
INSERT INTO login_attemp_repeatedly_stream;
```



© Paul Fremantle 2016. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. For more information, see <http://creativecommons.org/licenses/by-nc-sa/4.0/>

---

---

---

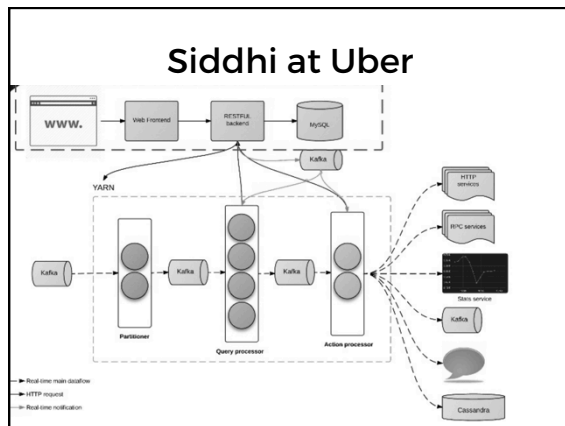
---

---

---

---

---




---

---

---

---

---

---

---

---

### Siddhi at Uber

- 100+ production apps
- 30 billion messages / day
- Fraud, anomaly detection
- Marketing, promotion
- Monitoring, feedback
- Real time analytics and visualization

<https://freo.me/siddhi-uber>

© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. <https://creativecommons.org/licenses/by-nc-sa/4.0/>

---

---

---

---

---

---

---

---

### Summary

- Realtime processing is hard
  - Requires large memory and state
  - The lambda architecture splits the problem into batch and realtime challenges
- Multiple approaches:
  - Pure Streaming
  - Micro-batch
  - CEP

© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. <https://creativecommons.org/licenses/by-nc-sa/4.0/>

---

---

---

---

---

---

---

---

## Questions?



© 2017 International 2017. This work is licensed under a Creative Commons  
Attribution-NonCommercial-ShareAlike 4.0 International License.  
For more information, see <https://creativecommons.org/licenses/by-nc-sa/4.0/>

---

---

---

---

---

---

---